

Recursive Feature Elimination by Sensitivity Testing

Nicholas Sean Escanilla*, Lisa Hellerstein[†], Ross Kleiman*, Charles Kuang*, James D. Shull[†] and David C. Page*

*Department of Computer Sciences

University of Wisconsin-Madison, Madison, Wisconsin

[†]Department of Oncology

University of Wisconsin-Madison, Madison, Wisconsin

[‡]Tandon School of Engineering

New York University, Brooklyn, New York

Abstract—There is great interest in methods to improve human insight into trained non-linear models, such as support vector machines (SVMs), deep neural networks, and large random forests; leading approaches include producing a ranking of the most relevant features, a non-trivial task for non-linear models. We show theoretically and empirically the benefit of a novel version of recursive feature elimination (RFE) as often used with SVMs; the key idea is a simple twist on the kinds of sensitivity testing employed in computational learning theory with membership queries (e.g., [1]). With membership queries, one can check whether changing the value of a feature in an example changes the label. In the real-world, we usually cannot get answers to such queries, so our approach instead makes these queries to a trained (imperfect) non-linear model. Because SVMs are widely used in bioinformatics, our empirical results use a real-world cancer genomics problem; because ground truth is not known for this task, we discuss the *potential* insights provided. We also evaluate on synthetic data where ground truth is known.

I. INTRODUCTION

There is great interest in methods to improve human insight into trained non-linear models such as support vector machines (SVMs), deep neural networks, and large random forests; one existing approach is to produce a ranking of the most relevant features, a non-trivial task for non-linear models. Famous examples of this approach include Breiman’s method for ranking features in a random forest by lost area under the Receiver Operating Characteristic curve (*AUC*) when a feature is deleted [2] and Guyon’s modified version of her recursive feature elimination (RFE) approach tailored to non-linear models [3], which both rank and performs feature selection. Breiman’s and Guyon’s approaches test how much the deletion of a feature costs in lost accuracy or *AUC*. In contrast, computational learning theory has a long history of *sensitivity* testing by “flipping” or changing the value of a feature, rather than deleting it, and posing a membership query to find the effect on the label of the example (e.g., [1]). In practice we do not have an oracle for such membership queries.

This paper presents an alternative algorithm, RFE by Sensitivity Testing (RFEST), that employs a trained non-linear model as an approximate oracle for such membership queries. Hence our algorithm asks how much accuracy or *AUC* is lost from a *trained* model when a variable is *flipped*, rather than how much is lost *compared to* an existing model when

a variable is *deleted*. We first prove a probably approximately correct (PAC)-like result showing that under certain assumptions this algorithm provides an accurate ranking; this result does not rely on any particular type of non-linear model or learning algorithm, but only on the condition that the algorithm achieves some minimum gain in accuracy over random guessing, as in weak learning. Second, we show empirically that when used with SVMs the new algorithm, RFEST, outperforms RFE in ranking (as a surrogate for insight) the genetic features associated with breast cancer in a genome-wide association study (GWAS) data set and on multiple synthetic data sets labeled by known ground truth, a family of arguably the most challenging non-linear target functions. Our study is limited to binary data but the algorithm potentially can be extended to other types of data by permuting feature values.

As a motivating example, genome-disease association studies (genome-wide or limited) seek genetic features associated with disease, i.e., predictive of disease. In many cases it is believed that such features may interact with one another in highly nonlinear ways to influence disease; nevertheless, for practical reasons almost all association studies use linear models and hence can find only features that individually are correlated with disease [4]. Consequently, key genetic features may be missed entirely. Because linear and non-linear SVMs have been widely used in bioinformatics applications, we will use SVMs as our learner for these empirical studies. The theoretical results show the algorithm can use any learner capable of building moderately accurate non-linear models.

II. BACKGROUND

A. Feature Ranking and Feature Selection

While feature ranking and feature selection are different problems, they are closely related and each is sometimes accomplished by the other. Recursive feature selection can rank by maintaining the order in which features are removed; feature ranking, e.g. by information gain, is often followed by removal of lower ranked features. Many feature selection algorithms utilize linear modeling approaches such as lasso-penalized logistic regression, linear SVMs, naïve Bayes or other weighted-voting schemes among features. An alternative is to implement these same approaches after filtering features

individually by information gain or by many single-variable logistic regression runs [5]. To account for interactions between features, the standard approach is to introduce interaction terms, but such terms typically are limited to pairs of features, and even then they greatly increase both run-time and risk of over-fitting.

Nonlinear SVMs have the potential to more effectively find complex interactions among features, but insight into the important interactions is hard to extract from the learned model. This paper addresses that shortcoming by presenting an alternative RFE algorithm and demonstrating that the algorithm makes it possible to identify the features that are relevant—that play a role in the learned nonlinear model, even if individually they are completely uncorrelated with the class—while removing those features that are irrelevant or redundant.

Because we do not assume we are in an active learning setting—we do not have access to an oracle for membership queries that can label feature vectors with the value of any feature altered—our key insight is to use the trained nonlinear SVM itself as such an oracle instead. While this trained model is not the target concept, we assume it is more accurate than random guessing and hence provides some information about feature relevance. In homage to earlier work on membership queries to test the sensitivity of target concepts to individual features, we call our RFE algorithm “RFE by Sensitivity Testing,” or RFEST. The remainder of the paper presents the RFEST algorithm and empirical evaluations of it, including novel and promising insights that it provides into genetic susceptibility to breast cancer.

B. SVMs, Correlation Immunity, and RFE

SVMs are known for both their strong performance and flexibility based on the chosen kernel [6]. The strength of SVMs comes from their ability to effectively learn nonlinear separators through use of the kernel trick, a mapping to a higher-dimensional feature space resulting in an ability to encode nonlinear separators in the original feature space [7].

Accordingly, it is expected that SVMs can efficiently learn correlation immune (CI) functions, which are notable nonlinear Boolean functions. A function is CI if every single-feature marginal distribution is uninformative, i.e., no feature by itself is correlated with the function value, or class, even given the entire truth table or example space. We say a function f is *correlation immune of order c* (or c -correlation immune) if f is statistically independent from any subset of variables with a size of at most c . A function is correlation immune if and only if every variable has zero gain (with respect to any gain measure) when computed from the input data (cf. [8]).

As a result, these functions include some of the most challenging target concepts for most classification algorithms, most noteworthy the parity functions. The most famous nonlinear separators in machine learning are exclusive-or (XOR) and exclusive-nor (XNOR), which are two-feature parity functions. These particular functions arise in practice, for example in biology (Table 1) [9]. In Table I, the interpretation of this

TABLE I
A TRUTH TABLE FOR DROSOPHILA (FRUITFLY) SURVIVAL BASED ON GENDER AND SXL GENE ACTIVITY.

GENDER FEMALE	SXL ACTIVE	SURVIVAL
0	0	0
0	1	1
1	0	1
1	1	0

output is that flies that will survive are either male with an active *Sxl* gene, or female with an inactive *Sxl* gene. While a nonlinear SVM can learn this function easily given only the relevant variables (i.e. *Gender Female* and *Sxl Active*), the SVM’s accuracy will degrade dramatically as irrelevant variables are added, unless the training set is quite large (see **Experimental Results**).

One would expect, for example, that SVMs, with a Gaussian kernel or polynomial kernel of degree at least two, would learn these functions with ease. Unfortunately, for the simplest case of XOR in the presence of even a modest number of irrelevant features, or variables, SVMs tend to have a difficult time learning and require a large sample size empirically. This problem is not specific to SVMs; it is known that no algorithm based on statistical queries can PAC learn parity functions of $\log(n)$ variables [10]. We seek a method of feature selection that can remove the irrelevant variables and restore classification performance.

A widely used approach to perform such a task is RFE, an embedded-based backward selection strategy [11]. RFE constructs an SVM, ranks the features according to the constructed SVM, removes the lowest ranked feature or features (e.g., bottom ten percent), and repeats until a certain (user-specified) number of features remain. The RFE algorithm with a linear SVM simply ranks features with respect to their given coefficients (i.e. from the learned model); this approach assumes features have been normalized to have comparable ranges.

Unfortunately, for a nonlinear SVM, feature coefficients cannot be obtained; Guyon *et al.* [3] presented a version of RFE for use with nonlinear SVMs. We propose an alternative RFE algorithm, RFEST, and compare these algorithms on synthetic data and a real-world cancer genomics problem.

C. Breast Cancer and Single-Nucleotide Polymorphisms

The development of breast cancer is influenced by many genetic and environmental factors. We study how feature selection performs on the variations at single base pairs of the human genome, which are known as single-nucleotide polymorphisms (SNPs). In cancer, both germline SNPs (the DNA sequence with which a person is born, and which is replicated in most of the cells in her body) and somatic mutations (variants that occur in select cells during replication and can lead to cancer) are important and are widely studied. To date, germline SNPs have received more attention as they can

predict a person's future risk of breast cancer [12]. Genome Wide Association Studies (GWAS) seek to find SNPs that are associated—correlated—with risk for developing disease.

Currently, GWAS consider SNPs independently and do not take into account possible interactions between SNPs. The rationale behind this is that it is infeasible, for example, to consider all pairs of the $n = 1$ million SNPs that are typically measured. The main purpose of a thorough investigation of SNPs is to gain a better understanding of how these genetic variants act as biological markers. Given a set of SNPs, if we can help identify a subset of important SNPs that correlate with a particular effect in patients, then we will be able to investigate their interactions. In turn, this will help our decision processes about numerous aspects of medical care such as the following: risk of developing a certain disease, effectiveness of various drugs, and adverse reactions to specific drugs.

III. ALGORITHMS

A. RFE Algorithm

In our experiments, we compare our RFEST algorithm to the RFE method proposed by Guyon *et al.* [3]. Although variants of RFE have been proposed [13]–[15], the original method of Guyon *et al.* is still widely used in the bioinformatics community [16]–[21]. Due to the nature of the data sets used in their paper, Guyon *et al.* utilized a linear SVM with RFE. However, they described how their method can be carried over to handle a nonlinear SVM implementation and this is the algorithm that we use as the baseline, which we describe next.

For SVMs, the cost function that is being minimized is the following:

$$J = \frac{1}{2} \alpha^T H \alpha - \alpha^T \mathbf{1} \quad (1)$$

with the following constraints:

$$0 \leq \alpha_k \leq C, \quad \sum_k \alpha_k y_k = 0$$

For training instances c and d , α is a vector of weights on the training instances learned by the SVM algorithm, y_c and y_d represent the class values for the training instances, \mathbf{x}_c and \mathbf{x}_d are the feature vectors for the training instances, C is a regularization parameter, and $H = y_c y_d K(\mathbf{x}_c, \mathbf{x}_d)$, for K a kernel function [3].

To determine feature relevance, the change in cost function proposed by Guyon is the following ranking coefficient:

$$DJ(j) = \frac{1}{2} \alpha^T H \alpha - \frac{1}{2} \alpha^T H(-j) \alpha \quad (2)$$

where $H(-j)$ represents a modified version of H that recomputes the matrix without the j^{th} feature. In turn, the feature with the smallest value for $DJ(j)$ is removed.

Algorithm 1 describes the RFE algorithm for the nonlinear case in more detail [22]. The benefit of using RFE over a vanilla approach (e.g. train a new SVM for each candidate feature on every iteration) allows for each iteration of the algorithm to train only one SVM model. In other words, we assume that the vector α is fixed and consider the change in the kernel as a result of removing feature j . Note that for

Algorithm 1 RFE Algorithm

Input: data $d_{i,j}$, where $i \in \{1, \dots, m\}, j \in \{1, \dots, n\}$
repeat
 Train SVM, output α
 Implement $DJ(j)$ according to (2), \forall features j
 Remove the feature(s) with the smallest $DJ(j)$
until k features remain ($k < n$)

each iteration, $H(-j)$ must be computed for each candidate feature j . The qualitative justification behind this cost function is that a feature's value to the learned model is measured by the change in the expected value of error when removing that candidate feature [22]. RFE iterates until k features remain, however to have a fair comparison to RFEST, the stopping criterion for RFE was set to be until the accuracy measurement AUC is less than the max AUC achieved thus far. We describe RFEST in the next section.

B. RFEST Algorithm

While our presentation of RFEST assumes that we are using binary features with a $\{-1, 1\}$ -encoding, it could be extended to handle continuous and/or categorical features as well. Standard RFE requires recomputing the H matrix (as described in the previous section) for each feature removed and can become computationally intractable with many features. In the case where there are thousands of features, Guyon *et al.* [3] chose to remove half of the features at each iteration. Doing so allows for faster convergence to an idealized subset of features, but key information may be lost.

There are two main differences between RFE and RFEST. The first is that RFEST flips the binary features, rather than deleting them. Note that flipping a feature means that if its current value is -1 , then it is changed to have the value 1 , and vice versa.

The second difference is in the construction of the cost function. An SVM classifier can classify a dataset with the accuracy measurement AUC . In addition, for each feature j , we create a modified version of the training set by flipping feature j in each example, and then calculate the AUC of the same SVM classifier on this modified training set. We call the calculated value $AUC_{flipped}$. The ranking coefficient used by RFEST is the following:

$$R(j) = AUC - AUC_{flipped} \quad (3)$$

The interpretation of $R(j)$ is as follows. For each j , if $AUC_{flipped} < AUC$, then the j^{th} feature is relevant because the model classified the instances at a lower AUC with j flipped. In contrast, if $AUC_{flipped} \geq AUC$, then the j^{th} feature is irrelevant because the model classified the instances with the same or higher AUC with j flipped. Therefore, the feature corresponding to the smallest $R(j)$ will be removed. This process does not retrain a classifier for every candidate feature to be removed and we no longer compute $H(-j)$.

For our experiments, a nonlinear SVM with an RBF kernel was used. The reason for doing so is because the RBF kernel

implicitly computes interaction terms for all subsets of input features. In addition, it has been shown that searching in exponentially growing sequences for the hyper-parameters, namely cost C and gamma γ , is a good method for identifying their respective parameter values [23]. Therefore, the best configuration for C and γ was chosen using grid search.

To determine the final subset of features, RFEST stops when the AUC at any given iteration is less than $p\%$ of the max AUC achieved thus far. This parameter controls the tradeoff between model interpretability and model efficacy. That is, a lower choice for p encourages a small number of coefficients, whereas a larger choice encourages a better performing model. For our experiments, we set $p = 95\%$ (see **Section IV**). Alternatively, other heuristics can be implemented with RFEST. One approach would be to stop when AUC decreases (i.e. a hill-climbing search). Another would be to apply a simulated annealing method, where if the AUC decreases, we continue searching with a small probability. As the search continues, the probability decreases. Search methods such as simulated annealing or the approach RFEST currently uses are aimed towards avoiding a local optimum. We use our current search method since it avoids having to set additional parameters.

RFEST may suffer if the SVM model we use is overfitted to a given instance, since the model might then be sensitive to the value of every feature. For this reason, we used a ten-fold cross validation and allocated the dataset into separate training, tuning, and testing sets to produce an unbiased estimate of the efficacy of our approach. The algorithm below summarizes RFEST.

Algorithm 2 RFEST Algorithm

Input: data $d_{i,j}$, where $i \in \{1, \dots, m\}$, $j \in \{1, \dots, n\}$
repeat
 Train SVM and output accuracy measurement AUC
 Implement $R(j)$ according to (3), \forall features j
 Remove the feature(s) with the smallest $R(j)$
until AUC is less than $p\%$ of the max AUC achieved

We next demonstrate the theoretical efficacy of using the quantity $R(j)$ to rank features, by considering a classification problem on n binary features, where examples are labeled according to the parity of a subset of the features. We show if a nonlinear machine learning algorithm can learn a sufficiently accurate model M , then with high probability, using a polynomial-size sample to compute the $R(j)$ values with model M will result in those values being higher for relevant features than for irrelevant features. Thus if any irrelevant features are present, the feature with the lowest $R(j)$ value, removed by RFEST, will be irrelevant.

RFEST is a generic method. Because of the interpretation behind the ranking coefficient $R(j)$, this allows us to use any accuracy measurement. Therefore, for simplicity and clarity of our analysis, we prove our theorem for a related measure, $\tilde{R}(j)$, which is the same as $R(j)$ except that it is defined in terms of accuracy rather than AUC . Although we state the theorem here only for the parity function and uniformly

distributed examples, we prove a more general theorem in the Supplementary Material. That theorem applies to a somewhat broader class of functions and product distributions.

Theorem III.1. *Let f be a Boolean target concept, defined on n Boolean features, which labels examples according to the parity of a fixed subset of the n features. Suppose a machine learning algorithm is used to learn a classifier M for f . Suppose further that M has true error rate $\epsilon < 1/2$, with respect to the uniform distribution. Then there is a quantity t that is polynomial in $n, \ln \frac{1}{\delta}$, and $\frac{1}{(1/2)-\epsilon}$, with the following property: for all $0 < \delta < 1$, if the $\tilde{R}(j)$ values for all n features are computed using M and a new independent sample of size t , drawn from the uniform distribution, then with probability at least $1 - \delta$, the computed $\tilde{R}(j)$ values for all the relevant features will be higher than the computed $\tilde{R}(j)$ values for the irrelevant features.*

While **Theorem III.1** may appear to contradict the known result that parity functions are not PAC-learnable from statistical queries, it does not because it is preconditioned on having an SVM model with accuracy better than random guessing.

IV. EXPERIMENTAL RESULTS

A. Data

We implemented RFEST and Guyon's RFE algorithm tailored to a nonlinear SVM, and we evaluated it on two types of data. The first consists of synthetic data that takes the form of a parity function on two variables, which is a CI function of order two. Correlation immune functions of order four, five, and six were also evaluated. There are many different CI functions, so for orders four, five, and six, ten functions for each order were randomly chosen. For functions of order c , the associated target concept was defined on n features. Of those n features, c were randomly chosen and corresponded to the c variables of the CI function, and the remaining features were irrelevant. Therefore, the task of both feature selection algorithms was to find the c variables that determined the class label. Feature values for all instances were chosen from a uniform distribution and the range of the number of instances was 100 to 2000.

The second dataset presented in this paper indicates that *Emca4*, a genetic determinant of susceptibility to 17β -estradiol (E2)-induced mammary cancer in the rat, has been mapped to rat chromosome 7 (RNO7) [24]–[26]. Data presented herein indicate that *Emca4* harbors multiple genetic determinants of mammary cancer susceptibility and tumor aggressiveness that are orthologous to breast cancer risk loci mapped to chromosome 8q24-24 in genome wide association studies (GWAS) [12], [27]–[31].

The proceeding algorithm(s) used 76 of the SNPs in the designated region that are in the Hunter GWAS data set of 1145 breast cancer cases and 1142 controls [32]. All patients that had incomplete SNP data (630 patients) were omitted from our analysis. The data was made available via dbGaP's Cancer Genetic Markers of Susceptibility (CGEMS) Breast Cancer GWAS.




Fig. 1. From left to right are results for 20 total features. The first column represents the AUC achieved across different training examples and the second column shows the number of features that were retained, with respect to the AUC achieved from the first column. The first row shows the results for RFE. The second row shows the results for RFEST.

B. Synthetic Data Results

The following were the different CI functions investigated: parity (order two), order four, order five, and order six. For orders four, five, and six, ten different functions per order were chosen for analysis. To best represent the results, a learning curve was created to show the average AUC for n total features, where $n \in \{20, 50, 100\}$ (i.e. the average AUC of the ten different functions for each order, respectively). For each n , we trained datasets that contained m examples, where $m \in \{100, 200, 300, \dots, 2000\}$. In addition, a learning curve was plotted to represent the average number of features that were kept using the same number of features and examples as stated above (i.e. the average number of features retained of the ten different functions for each order, respectively). To make the comparison fair, 10% percent of the total number of features remaining at a given iteration were removed. In [3], the authors removed half of the features. However, we believe removing 10% of features at each iteration gave more accurate results since the number of features to begin with was not as large as the number in Guyon's paper.

In Figure 1 we show the synthetic data results for 20 total features across the CI functions of order two, four, five, and six. A learning curve was plotted to show the results of the average AUC and the average number of features retained across the various orders. The figures show the outcome for the RFEST and RFE algorithm. The performance (in terms of AUC) of both algorithms increased as the number of training examples increased, which is to be expected. However, RFEST achieved an AUC of 1.0 at a faster rate than RFE. In fact, for orders five and six, RFE failed to attain an AUC of 1.0.

The average number of features retained across the various orders was also calculated (second column of Figure 1). The

goal was to output only the relevant features. For example, in the case of the parity function, we set the relevant features to be randomly chosen among the 20 features in our dataset, and the remaining features were irrelevant. For the CI function of order four, four randomly chosen features were set as the relevant variables, and the remaining features were irrelevant. The creation of the remaining CI functions followed a similar format. All feature values were chosen with respect to a uniform distribution.

Observe that in Figure 1, RFE was not able to retain the relevant features across all orders. For orders four, five, and six, the algorithm stopped prematurely, outputting nearly all of the original features. In the case of the parity function, there are several instances where the RFE algorithm outputted only a small subset of features that included the relevant variables, however, it failed to return solely those that are relevant. Unlike RFE, RFEST was able to retain solely the relevant variables for each order. For the parity function, only 200 instances were required. For orders four, five, and six, 400, 400, and 300 instances were needed to return a subset of only the relevant features, respectively. This is a significant difference.

Figure 2 shows the synthetic data results for 50 total features across the CI functions of order two, four, five, and six. In a similar format to Figure 1, the first row represents the AUC achieved and features retained, respectively, for RFE. The second row represents the results for RFEST. As compared to Figure 1, there is a general decrease in AUC across all orders, as the number of irrelevant features increases. However, after a certain number of training examples, RFEST outperformed RFE. The max average AUC achieved for RFE and RFEST for all orders are represented in Table II.

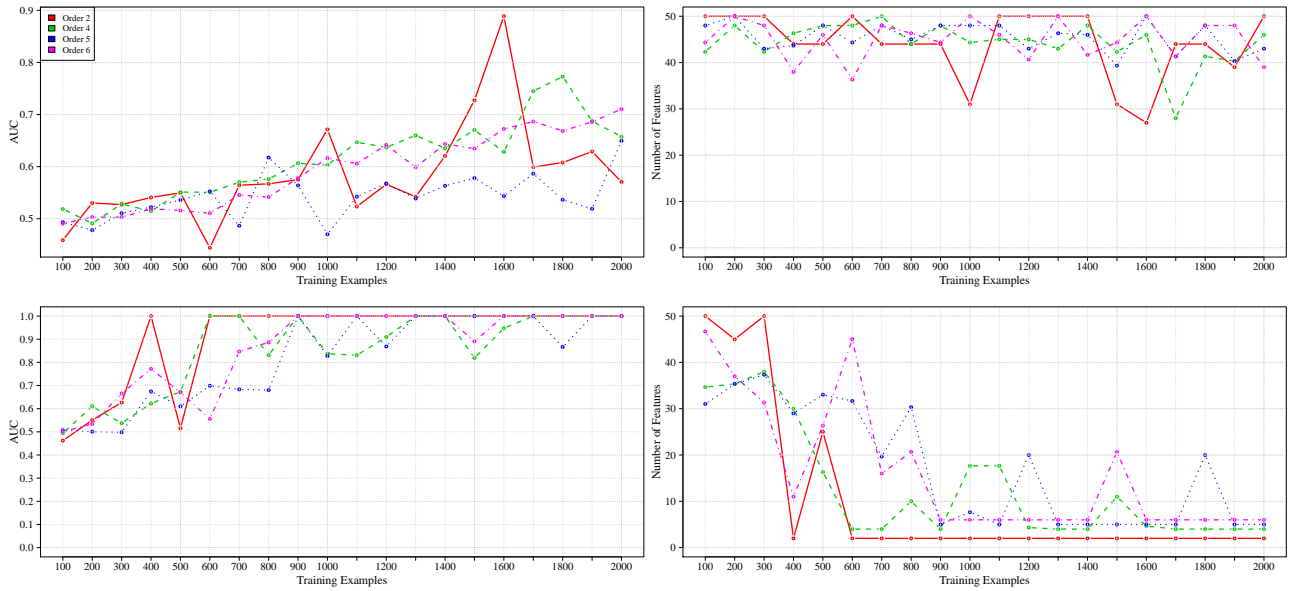


Fig. 2. From left to right are results for 50 total features. The first column represents the AUC achieved across different training examples and the second column shows the number of features that were retained, with respect to the AUC achieved from the first column. The first row shows the results for RFE. The second row shows the results for RFEST.

In addition to the significant difference in performance (as shown in Table II), there is a distinct difference in the number of features returned. Across all orders and all varying training examples, RFE was not able to find the subset of relevant variables. However, RFEST was able to do so with the corresponding max average AUC 's and training examples from Table II. This observation is also shown graphically in the second column of Figure 2.

TABLE II
MAX AVERAGE AUC RESULTS FOR 50 TOTAL FEATURES.

ORDER	RFE		RFEST	
	AUC	TRAINING EXAMPLES	AUC	TRAINING EXAMPLES
2	0.889	1600	1.0	400
4	0.773	1800	1.0	600
5	0.649	2000	1.0	900
6	0.710	2000	1.0	900

Lastly, Figure 3 shows the synthetic data results for 100 total features across the CI functions of order two, four, five, and six. Similar to the results in Figures 1 and 2, RFEST outperformed RFE in both prediction performance and the ability to retain fewer relevant features. The max average AUC results for RFE and RFEST can be found in Table III. For 100 total features, at approximately 900 training examples, there is a significant difference between the prediction performance (across CI function orders two, four, and six) for RFE and RFEST (as shown in Table III). That is, with fewer training examples, RFEST achieved higher max average AUC 's compared to RFE.

Observe that in Figure 3, as the number of training examples increased, RFEST was able to return a smaller subset of features, whereas, RFE returned more than half the number of original features. Note that RFEST demonstrated more fluctuation with 100 total features. This finding suggests that there are CI functions in which RFEST may not be the more robust method, in terms of its ability to return relevant features.

TABLE III
MAX AVERAGE AUC RESULTS FOR 100 TOTAL FEATURES.

ORDER	RFE		RFEST	
	AUC	TRAINING EXAMPLES	AUC	TRAINING EXAMPLES
2	0.557	1100	1.0	900
4	0.624	1700	1.0	1600
5	0.548	1800	0.704	1500
6	0.611	1800	1.0	1400

C. Germline Genomic Data for Breast Cancer Results

There is great interest in associating variations in the human genome with disease risk. Much of this work focuses on associating with any given disease the variations in SNPs. Most such work assumes the SNPs, and the variations in disease risk that they cause, are independent of one another; in general this assumption is wrong and results in lost accuracy.

We may examine variations in the germline DNA with which a person is born or variations that arise from somatic mutations in individual cells, such as in the development of cancers and which may vary widely even within the same tumor. One particular disease for which both types of

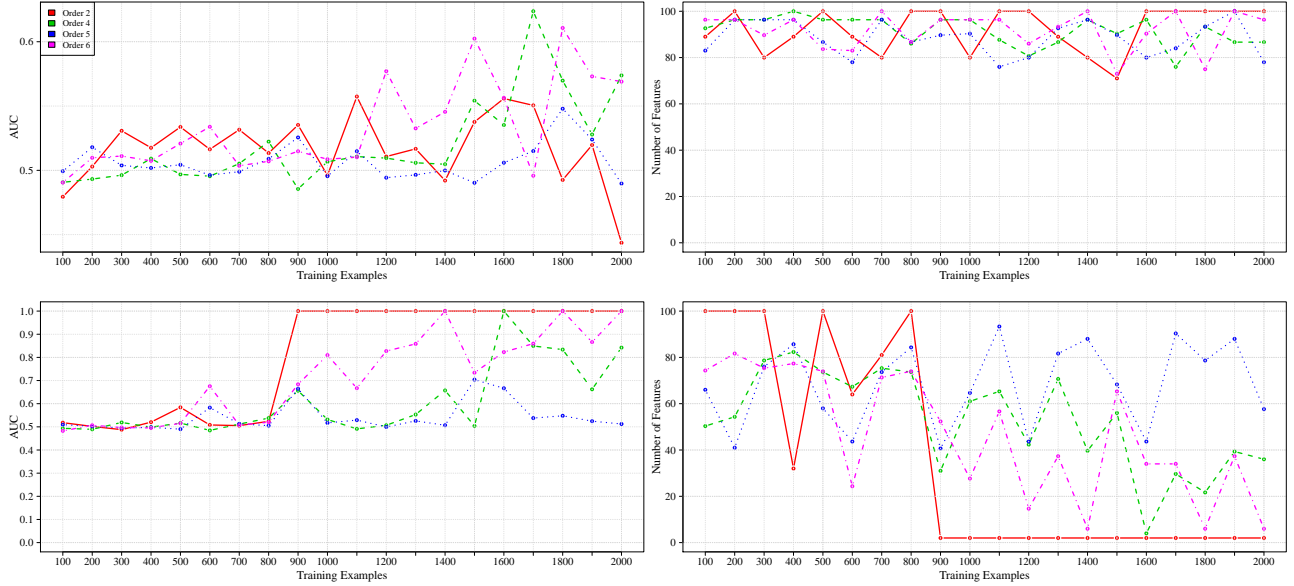


Fig. 3. From left to right are results for 100 total features. The first column represents the AUC achieved across different training examples and the second column shows the number of features that were retained, with respect to the AUC achieved from the first column. The first row shows the results for RFE. The second row shows the results for RFEST.

variations have been studied is breast cancer. For predicting disease risk, germline genomic data is the more natural choice to use.

The data we investigate contains 76 SNPs, translating to 152 binary features, in a particular region of the human genome that is orthologous to a region of the rat genome known to modulate breast cancer risk. We use the CGEMS data set of SNP genotypes for 1145 breast cancer cases and 1142 healthy age- and gender-matched controls [32]. Applying RFEST to this data, as run in the previous section, produces a cross-validated AUC of 0.56 with only nine features, which outperforms linear SVM and nonlinear SVM cross-validated runs with the original input data (0.53 and 0.54, respectively). Likewise, RFEST outperformed RFE as RFE (as run in the previous section) returned an AUC of 0.53 with 122 features, no better than either a linear or nonlinear SVM run. We also implemented RFE with the stopping criterion stated in Algorithm 1. That is, since RFEST returned nine features, RFE also iterated until nine features remained but returned an AUC of 0.51.

While all runs were performed by eliminating 10% of features at a time, our novel algorithm is also effective (when compared to RFE) when removing 20% or even 30% of the remaining features at a time. Removing 10% of the features at a time not only resulted in an AUC of 0.56 but most notably retained only nine features. With such a small set of features, one can then exhaustively generate all pairs (and even more) of interaction terms. A linear SVM model was built with the remaining nine features and all interaction terms. In turn, the top 13 features were all pairs of SNPs rather than individual SNPs. This suggests that interactions play a major role in the effect of SNP variations in this region on breast cancer risk, as

has been suspected. Studies are under way to further evaluate these nine selected SNPs.

It has been shown that incorporating, as risk factors, germline SNPs associated with breast cancer can significantly improve prediction and even mammography-based diagnosis of breast cancer, even though breast cancer is estimated to be only 30% heritable [33]. In this section we show that avoiding the independence assumption regarding SNPs, by using a nonlinear SVM with our novel RFE algorithm, makes it possible to associate with breast cancer new SNPs and their interactions, and that this association can enable more accurate breast cancer risk prediction than could be made from these SNPs without taking interactions into account.

A post-hoc analysis of these nine selected SNPs confirmed our collaborating biologist's suspicion that interactions (rather than specific SNP values) were the most important modulator of breast cancer risk in this genomic region, and revealed which interactions were crucial to the underlying task (i.e. breast cancer diagnoses).

V. DISCUSSION & CONCLUSION

In this paper, we explored the difficulties that accompany feature selection and learning decidedly nonlinear target concepts. In addition, we discussed the challenges that are faced in using Guyon's RFE algorithm [3]. Such problems occur in runtime for large datasets as well as a significantly lower AUC than the novel RFE algorithm.

We introduce a new algorithm, RFEST, for a nonlinear machine learning algorithm and demonstrate its efficacy both theoretically (refer to **Theorem III.1** and Supplementary Material) and empirically (see **Section IV**). The RFE algorithm is an embedded-based approach but RFEST behaves like a

wrapper-based approach. It uses a nonlinear SVM as a black box to determine feature relevance. In principle, with this approach one can use *any* machine learning algorithm to remove irrelevant or redundant features.

RFEST differs from RFE in two important ways: it perturbs rather than eliminates each feature to test sensitivity and measures loss in model efficacy instead of the loss in weighted sum of distances from the margin. These differences result in substantial improvements across CI functions and a real-world breast cancer genomics problem.

Extending the feature types used by RFEST is left for future work. Lastly, if one knew that the input data contained many correlated features, then applying a filter algorithm before RFEST will aid in removing redundant features.

ACKNOWLEDGMENT

The authors would like to thank...

REFERENCES

- [1] N. Bshouty, T. Hancock, and L. Hellerstein, "Learning boolean read-once formulas with arbitrary symmetric and constant fan-in gates," in *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, ser. COLT '92. New York, NY, USA: ACM, 1992, pp. 1–15.
- [2] L. Breiman, "Random forests," *Mach. Learn.*, vol. 45, no. 1, pp. 5–32, Oct. 2001.
- [3] I. Guyon, J. Weston, S. Barnhill, and V. Vapnik, "Gene selection for cancer classification using support vector machines," *Machine Learning*, vol. 46, no. 1, pp. 389–422, 2002.
- [4] Z. Zhang, E. Ersoz, C. Lai, R. Todhunter, H. Tiwari, M. Gore, ..., and E. Buckler, "Mixed linear model approach adapted for genome-wide association studies," *Nature*, vol. 42, pp. 355 EP –, 2010.
- [5] G. Chandrashekar and F. Sahin, "A survey on feature selection methods," *Computers and Electrical Engineering*, vol. 40, no. 1, pp. 16 – 28, 2014, 40th-year commemorative issue.
- [6] R. Caruana and A. Niculescu-Mizil, "An empirical comparison of supervised learning algorithms," in *Proceedings of the 23rd International Conference on Machine Learning*, ser. ICML '06. New York, NY, USA: ACM, 2006, pp. 161–168.
- [7] S. Maldonado and R. Weber, *Embedded Feature Selection for Support Vector Machines: State-of-the-Art and Future Challenges*. Springer Berlin Heidelberg, 2011.
- [8] B. Roy, "A brief outline of research on correlation immune functions," in *Information Security and Privacy*, L. Batten and J. Seberry, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2002, pp. 379–394.
- [9] T. Cline, "A male-specific lethal mutation in drosophila melanogaster that transforms sex," *Developmental Biology*, vol. 72, pp. 266–275, 1979.
- [10] A. Blum, M. L. Furst, J. C. Jackson, M. J. Kearns, Y. Mansour, and S. Rudich, "Weakly learning DNF and characterizing statistical query learning using fourier analysis," in *Proceedings of the Twenty-Sixth Annual ACM Symposium on Theory of Computing*, 23-25 May 1994, Montréal, Québec, Canada, 1994, pp. 253–262.
- [11] C. Stambaugh, H. Yang, and F. Breuer, "Analytic feature selection for support vector machines," *CoRR*, vol. abs/1304.5678, 2013.
- [12] K. Michailidou, S. Lindström, J. Dennis, J. Beesley, S. Hui, S. Kar, ..., and D. Easton, "Association analysis identifies 65 new breast cancer risk loci," *Nature*, vol. 551, pp. 92 EP –, 2017.
- [13] M. Nguyen and F. de la Torre, "Optimal feature selection for support vector machines," *Pattern Recognition*, vol. 43, no. 3, pp. 584 – 591, 2010.
- [14] T. Wang, H. Huang, S. Tian, and J. Xu, "Feature selection for svm via optimization of kernel polarization with gaussian ard kernels," *Expert Syst. Appl.*, vol. 37, no. 9, pp. 6663–6668, Sep. 2010.
- [15] Q. Liu, C. Chen, Y. Zhang, and Z. Hu, "Feature selection for support vector machines with rbf kernel," *Artificial Intelligence Review*, vol. 36, no. 2, pp. 99–115, Aug 2011.
- [16] P. Tao, T. Liu, X. Li, and L. Chen, "Prediction of protein structural class using tri-gram probabilities of position-specific scoring matrix and recursive feature elimination," *Amino Acids*, vol. 47, no. 3, pp. 461–468, Mar 2015.
- [17] M. Schwartz, Z. Hou, N. Propson, J. Zhang, C. Engstrom, V. Costa, ..., and J. Thomson, "Human pluripotent stem cell-derived neural constructs for predicting neural toxicity," *Proceedings of the National Academy of Sciences*, vol. 112, no. 40, pp. 12 516–12 521, 2015.
- [18] M. Qureshi, B. Min, H. Jo, and B. Lee, "Multiclass classification for the differential diagnosis on the adhd subtypes using recursive feature elimination and hierarchical extreme learning machine: Structural mri study," *PLOS ONE*, vol. 11, no. 8, pp. 1–20, 08 2016.
- [19] E. Zarogianni, A. Storkey, E. Johnstone, D. Owens, and S. Lawrie, "Improved individualized prediction of schizophrenia in subjects at familial high risk, based on neuroanatomical data, schizotypal and neurocognitive features," *Schizophrenia Research*, vol. 181, pp. 6 – 12, 2017.
- [20] B. Kampe, S. Klotz, T. Bocklitz, P. Rösch, and J. Popp, "Recursive feature elimination in raman spectra with support vector machines," *Frontiers of Optoelectronics*, vol. 10, no. 3, pp. 273–279, Sep 2017.
- [21] L. Kong, L. Kong, C. Wang, R. Jing, and L. Zhang, "Predicting protein structural class for low-similarity sequences via novel evolutionary modes of pseac and recursive feature elimination," *Letters in Organic Chemistry*, vol. 14, no. 9, pp. 673–683, 2017.
- [22] T. Lal, O. Chapelle, J. Weston, and A. Elisseeff, *Embedded methods*, ser. Studies in Fuzziness and Soft Computing ; 207. Berlin, Germany: Springer, 2006, pp. 137–165.
- [23] C. Hsu, C. Chang, and C. Lin, "A practical guide to support vector classification," 2010.
- [24] J. Colletti, K. Leland-Wavrin, S. Kurz, M. Hickman, N. Seiler, N. Samanas, ..., and J. Shull, "Validation of six genetic determinants of susceptibility to estrogen-induced mammary cancer in the rat and assessment of their relevance to breast cancer risk in humans," *G3*, vol. 4, no. 8, pp. 1385–1394, 2014.
- [25] B. Schaffer, C. Lachel, K. Pennington, C. Murrin, T. Strecker, M. Tochacek, ..., and J. Shull, "Genetic bases of estrogen-induced tumorigenesis in the rat: Mapping of loci controlling susceptibility to mammary cancer in a brown norway x aci intercross," *Cancer Research*, vol. 66, no. 15, pp. 7793–7800, 2006.
- [26] J. Shull, "The rat oncogenome: Comparative genetics and genomics of rat models of mammary carcinogenesis," *Breast Disease*, vol. 28, no. 1, pp. 69–86, 2007.
- [27] D. Easton, K. Pooley, A. Dunning, P. Pharoah, D. Thompson, D. B. ..., and B. Ponder, "Genome-wide association study identifies novel breast cancer susceptibility loci," *Nature*, vol. 447, no. 7148, pp. 1087–1093, 2007.
- [28] C. Turnbull, S. Ahmed, J. Morrison, D. Pernet, A. Renwick, M. Maranian, ..., and D. Easton, "Genome-wide association study identifies five new breast cancer susceptibility loci," *Nat Genet*, vol. 42, no. 6, pp. 504–507, 2010.
- [29] O. Fletcher, N. Johnson, N. Orr, F. Hosking, L. Gibson, K. Walker, ..., and J. Peto, "Novel breast cancer susceptibility locus at 9q31.2: Results of a genome-wide association study," *Journal of the National Cancer Institute*, vol. 103, pp. 425–435, 2011.
- [30] K. Michailidou, P. Hall, A. Gonzalez-Neira, M. Ghoussaini, J. Dennis, R. Milne, ..., and D. Easton, "Large-scale genotyping identifies 41 new loci associated with breast cancer risk," *Nat Genet*, vol. 45, no. 4, pp. 353–361, 2013.
- [31] H. Ahsan, J. Halpern, M. Kibriya, B. Pierce, L. Tong, E. Gamazon, ..., and A. Whittemore, "A genome-wide association study of early-onset breast cancer identifies ptkm as a novel breast cancer gene and supports a common genetic spectrum for breast cancer at any age. cancer epidemiology, biomarkers and prevention : a publication of the american association for cancer research, cosponsored by the american society of preventive oncology," *Cancer Epidemiology and Prevention Biomarkers*, vol. 23, no. 4, pp. 658–669, 2014.
- [32] D. Hunter, P. Kraft, K. Jacobs, D. Cox, M. Yeager, S. Hankinson, ..., and S. Chanock, "A genome-wide association study identifies alleles in fgfr2 associated with risk of sporadic postmenopausal breast cancer," *Nat Genet*, vol. 39, no. 7, pp. 870–874, 2007.
- [33] J. Liu, D. Page, P. Peissig, C. McCarty, A. Onitilo, A. Trentham-Dietz, and E. Burnside, "New genetic variants improve personalized breast cancer diagnosis," *AMIA Jt Summits Transl Sci Proc*, pp. 83–89, 2014.